

08

- **Open-model ecosystem** (Qwen, Apple MLX, Ollama): the open weights and runtimes both judges run on
- **Security resellers & MSSPs** who bundle Senti into their managed services
- **Tumo, universities & startup programs** for first pilots and early users
- **Open threat data** (prompt-injection datasets, attack research) for rules and benchmarks

- **Security advisor** (current or former CISO) who opens doors to buyers, for a small equity stake
- **Open-source contributors** who build adapters and rules, cutting our upkeep cost
- **Alumni network** (Tumo, universities) at software companies: warm intros to CTOs and security leads

07

- Keep adapters working as 7 agents ship updates
- Tune rules, detectors and the judge; red-team new attack patterns
- Build role-based policy profiles
- Run pilots and pass customer security reviews

06

- Fail-closed engine: rules → profile → detectors → judge
- On-device judge (Qwen3-4B on MLX)
- Hash-chained audit, undo snapshots, honeytokens, secret broker
- Agent-action benchmarks from the prototype

02

Role profiles set once in the admin panel, enforced the same way in all 7 agents. Hard rules no AI judge can overrule.

Every action is checked before it runs. If one slips through: file undo, a tamper-evident audit and decoy secrets. Real keys never reach the agent.

Judging runs on the Mac or your own server.
Senti never sees your data, so the vendor
review stays short.

Quiet by default: safe actions pass silently, so developers don't work around it.

Works with Claude Code · Codex · Cursor ·
Cline · OpenCode · Hermes · OpenClaw

04

- **Admin panel** as the daily touchpoint: live view of agent activity
- **Guided pilots:** org stack runs with one docker compose command
- **Hands-on onboarding** and policy tuning for large accounts
- Ongoing rule and threat updates

03

- **Direct sales** to CISOs and heads of IT at Mac-heavy companies
- **Content for security teams:** write-ups of real agent attacks and how to stop them
- **Live demo:** a poisoned-repo attack blocked in real time
- MDM marketplaces, resellers, security meetups

01

- **Security & IT teams** at companies whose staff run AI agents on Macs

- Mid-size software companies and agencies with Mac-heavy teams
- **Regulated industries** (fintech, health, legal) that can't send activity data to a cloud service

- The customer's employees (developers, knowledge workers) running agents every day. They don't pay, but they must not feel slowed down

09

- **Engineering:** engine, admin panel and 7 agent adapters that need constant upkeep
- **Sales & trust:** pilots, security reviews, certifications (e.g. SOC 2)
- **Model R&D:** judge evaluation, benchmarks, red-teaming
- **Low serving cost:** AI runs on customers' Macs and servers, not our cloud, so margins stay high

05

- **Per-device subscription:** admin panel, role profiles, corporate judge, audit. \$100 per Mac / year (about \$8 a month)
- **Paid pilots** and onboarding services
- **Enterprise add-ons (planned):** SSO, SIEM export, priority support
- **Annual contracts** per company, with volume pricing for large fleets